

This document, "Lecture 41: Multilayer Feed Forward Neural Network," discusses Multilayer Perceptron Neural Networks (MLP NNs), their architecture, and applications.

Key points include:

- **Multilayer Perceptron Neural Network:**
  - No connections within a layer or direct connections between input and output layers.
  - Fully connected between layers, often with more than three layers.
  - The number of output units, and hidden units per layer, can differ from the number of input units.
  - Each unit functions as a perceptron.
- **Backpropagation:** Details the mathematical calculations for activation, output, error, and weight updates between hidden and output units.
- **Applications:**
  - Digit Recognition (e.g., using 16 input values and 4 output bits to classify digits 0-9).
  - General applications include pattern classification, image processing, speech analysis, optimization problems, stock market forecasting, simulation, and robot steering.
- **Design Issues:** Covers the number of nodes in input and output layers, network topology (number of hidden layers and nodes, feed-forward or recurrent), weight and bias initialization, and handling missing values in training examples.
- **Characteristics of ANNs:**
  - Multilayer NNs with at least one hidden layer are universal approximators.
  - They can handle redundant features by learning small weights for them.
  - NNs are sensitive to noise.
  - Training can be time-consuming, but classification of test examples is rapid once training is complete.
  - Only network information needs to be stored for testing, reducing space requirements.

## Lecture 42

This file, "lecture42.pdf," covers various topics related to hypothesis learning, class imbalance, and performance metrics in classification models.

Here's a summary of its key sections:

- **Hypothesis Learning:**
  - Uses logical inference to find a hypothesis that matches examples.
  - Hypotheses are logical sentences that classify examples.
  - The aim of inductive learning is to find a hypothesis that classifies well and generalizes to new examples.
  - Discusses how hypotheses inconsistent with new examples are removed.
  - Introduces "false negatives" (hypothesis says negative, actually positive) and "false positives" (hypothesis says positive, actually negative).
  - The "Current-Best-Learning" algorithm is presented, which searches for a consistent hypothesis and backtracks when needed.
- **Version Space Learning (Candidate Elimination Algorithm):**
  - An alternative to backtracking, it keeps all hypotheses consistent with the data so

- far.
- Each new example removes inconsistent hypotheses, shrinking the "version space."
- This approach is incremental, meaning old examples don't need re-examination.
- The challenge of representing an enormous hypothesis space is addressed by "Least Commitment Search (LCS)."
- **Least Commitment Search (LCS):**
  - Uses a partial ordering (generalization/specialization) on the hypothesis space.
  - The entire version space is represented by two boundary sets: a "most general boundary" (G-set) and a "most specific boundary" (S-set).
  - All hypotheses between the S-set and G-set are consistent with the examples.
  - Describes how to initialize (G-set to True, S-set to False) and update S and G sets when new examples arrive, handling false positives and false negatives.
  - Discusses three possible outcomes: a single hypothesis remains, the version space collapses (no consistent hypothesis), or multiple hypotheses remain (representing a disjunction).
- **Drawbacks to LCS:**
  - Susceptible to collapsing if the domain contains noise or insufficient attributes.
  - Can lead to exponentially growing S-set or G-set elements in some hypothesis spaces.
- **Class Imbalance Problem:**
  - Occurs when one class (majority class) is significantly more represented than others (minority class).
  - Standard classifiers tend to be biased towards the majority class, leading to poor performance on the minority class.
  - Examples include fraud detection, anomaly detection, and medical diagnosis.
- **Approaches to Address Class Imbalance:**
  - **Data Level (Re-Sampling):**
    - **Oversampling:** Increases minority class instances (e.g., simple duplication, SMOTE). Can improve performance but risks overfitting. SMOTE creates synthetic samples.
    - **Undersampling:** Decreases majority class instances (e.g., random undersampling, Tomek Links). Can reduce training time but risks losing important data.
  - **Algorithmic Level:** Cost-Sensitive Learning, Ensemble Methods, Focal Loss.
  - **Other Approaches:** Batch Size Considerations, Feature Selection.
- **Key Considerations for Resampling:**
  - Severity of imbalance, dataset size, specific model, and risk of overfitting.
- **Performance Metrics:**
  - **Confusion Matrix:** A table describing a classification model's performance, showing true positives (TP), true negatives (TN), false positives (FP), and false negatives (FN).
  - **Precision:** The ability to retrieve mostly relevant documents ( $P = TP / (TP + FP)$ ).
  - **Recall:** The ability to find all relevant items ( $R = TP / (TP + FN)$ ).
  - **F-Measure (F1/Harmonic Mean):** Combines precision and recall, useful for skewed data where accuracy is not appropriate ( $F = 2RP / (R + P)$ ).
  - **Accuracy:** Overall correctness (Accuracy =  $(TP + TN) / (TP + FP + FN + TN)$ ).

- Emphasizes that accuracy alone is insufficient for imbalanced datasets.
- **ROC Curves:** Visualize a classifier's performance across different thresholds, illustrating the trade-off between True Positive Rate (TPR) and False Positive Rate (FPR). Higher curves indicate better performance.
- **Cost-Sensitive Learning:**
  - Aims to minimize the expected cost of misclassifications, as different errors can have different costs (e.g., medical diagnosis, fraud detection).
  - Illustrates with a "Cost Matrix" how a model with lower accuracy might be preferred if it minimizes overall cost.